

Additional callbacks

July 30, 2021

1 Additional callbacks

In this reading we'll be looking at more of the inbuilt callbacks available in Keras.

```
In [1]: import tensorflow as tf
        print(tf.__version__)
```

2.0.0

We will again be using the sklearn diabetes dataset to demonstrate these callbacks.

```
In [2]: # Load the diabetes dataset
```

```
        from sklearn.datasets import load_diabetes
```

```
        diabetes_dataset = load_diabetes()
```

```
In [3]: # Save the input and target variables
```

```
        from sklearn.model_selection import train_test_split
```

```
        data = diabetes_dataset['data']
```

```
        targets = diabetes_dataset['target']
```

```
In [4]: # Split the data set into training and test sets
```

```
        train_data, test_data, train_targets, test_targets = train_test_split(data, targets, t
```

Let's also build a simple model to fit to the data with our callbacks.

```
In [5]: # Build the model
```

```
        from tensorflow.keras.models import Sequential
```

```
        from tensorflow.keras.layers import Dense
```

```
        model = tf.keras.Sequential([
```

```
            Dense(128, activation='relu', input_shape=(train_data.shape[1],)),
```

```

        Dense(64,activation='relu'),
        Dense(64, activation='relu'),
        Dense(64, activation='relu'),
        Dense(1)
    ])

```

In [6]: *# Compile the model*

```

model.compile(loss='mse',
              optimizer="adam",metrics=["mse","mae"])

```

Now onto the callbacks!

1.1 Learning rate scheduler

Usage: `tf.keras.callbacks.LearningRateScheduler(schedule, verbose=0)`

The learning rate scheduler that we implemented in the previous reading as a custom callback is also available as a built in callback.

As in our custom callback, the `LearningRateScheduler` in Keras takes a function `schedule` as an argument.

This function `schedule` should take two arguments: * The current epoch (as an integer), and * The current learning rate, and return new learning rate for that epoch.

The `LearningRateScheduler` also has an optional `verbose` argument, which prints information about the learning rate if it is set to 1.

Let's see a simple example.

In [7]: *# Define the learning rate schedule function*

```

def lr_function(epoch, lr):
    if epoch % 2 == 0:
        return lr
    else:
        return lr + epoch/1000

```

In [8]: *# Train the model*

```

history = model.fit(train_data, train_targets, epochs=10,
                  callbacks=[tf.keras.callbacks.LearningRateScheduler(lr_function, v

```

Epoch 00001: LearningRateScheduler reducing learning rate to 0.0010000000474974513.

Epoch 00002: LearningRateScheduler reducing learning rate to 0.0020000000474974513.

Epoch 00003: LearningRateScheduler reducing learning rate to 0.0020000000949949026.

Epoch 00004: LearningRateScheduler reducing learning rate to 0.005000000094994903.

```
Epoch 00005: LearningRateScheduler reducing learning rate to 0.004999999888241291.
Epoch 00006: LearningRateScheduler reducing learning rate to 0.009999999888241292.
Epoch 00007: LearningRateScheduler reducing learning rate to 0.009999999776482582.
Epoch 00008: LearningRateScheduler reducing learning rate to 0.01699999977648258.
Epoch 00009: LearningRateScheduler reducing learning rate to 0.016999999061226845.
Epoch 00010: LearningRateScheduler reducing learning rate to 0.025999999061226846.
```

You can also use lambda functions to define your schedule given an epoch.

In [9]: *# Train the model with a difference schedule*

```
history = model.fit(train_data, train_targets, epochs=10,
                    callbacks=[tf.keras.callbacks.LearningRateScheduler(lambda x:1/(3+
verbose=False)
```

```
Epoch 00001: LearningRateScheduler reducing learning rate to 0.3333333333333333.
Epoch 00002: LearningRateScheduler reducing learning rate to 0.125.
Epoch 00003: LearningRateScheduler reducing learning rate to 0.07692307692307693.
Epoch 00004: LearningRateScheduler reducing learning rate to 0.05555555555555555.
Epoch 00005: LearningRateScheduler reducing learning rate to 0.043478260869565216.
Epoch 00006: LearningRateScheduler reducing learning rate to 0.03571428571428571.
Epoch 00007: LearningRateScheduler reducing learning rate to 0.030303030303030304.
Epoch 00008: LearningRateScheduler reducing learning rate to 0.02631578947368421.
Epoch 00009: LearningRateScheduler reducing learning rate to 0.023255813953488372.
Epoch 00010: LearningRateScheduler reducing learning rate to 0.020833333333333332.
```

1.2 CSV logger

Usage `tf.keras.callbacks.CSVLogger(filename, separator=',', append=False)`

This callback streams the results from each epoch into a CSV file. The first line of the CSV file will be the names of pieces of information recorded on each subsequent line, beginning with the epoch and loss value. The values of metrics at the end of each epoch will also be recorded.

The only compulsory argument is the filename for the log to be streamed to. This could also be a filepath.

You can also specify the separator to be used between entries on each line.

The append argument allows you the option to append your results to an existing file with the same name. This can be particularly useful if you are continuing training.

Let's see an example.

```
In [10]: # Train the model with a CSV logger
```

```
history = model.fit(train_data, train_targets, epochs=10,
                    callbacks=[tf.keras.callbacks.CSVLogger("results.csv")], verbose=1)
```

Let's view the information in the CSV file we have created using pandas.

```
In [11]: # Load the CSV
```

```
import pandas as pd

pd.read_csv("results.csv", index_col='epoch')
```

```
Out[11]:
```

	loss	mae	mse
epoch			
0	5808.937287	65.01146	5808.9370
1	5803.596608	64.81101	5803.5960
2	5814.807540	65.03268	5814.8080
3	5806.244453	64.88313	5806.2440
4	5804.998943	64.85107	5804.9995
5	5808.092012	64.83402	5808.0920
6	5821.321205	65.28305	5821.3213
7	5801.435458	65.08316	5801.4355
8	5810.194807	64.96325	5810.1953
9	5819.899328	65.07032	5819.8994

1.3 Lambda callbacks

Usage `tf.keras.callbacks.LambdaCallback(on_epoch_begin=None, on_epoch_end=None, on_batch_begin=None, on_batch_end=None, on_train_begin=None, on_train_end=None)`

Lambda callbacks are used to quickly define simple custom callbacks with the use of lambda functions.

Each of the functions require some positional arguments. *on_epoch_begin and on_epoch_end expect two arguments: epoch and logs, *on_batch_begin and on_batch_end expect two arguments: batch and logs and *on_train_begin and on_train_end expect one argument: logs.

Let's see an example of this in practice.

```
In [12]: # Print the epoch number at the beginning of each epoch
```

```
epoch_callback = tf.keras.callbacks.LambdaCallback(
    on_epoch_begin=lambda epoch,logs: print('Starting Epoch {}'.format(epoch+1)))
```

```
In [13]: # Print the loss at the end of each batch
```

```
batch_loss_callback = tf.keras.callbacks.LambdaCallback(  
    on_batch_end=lambda batch,logs: print('\n After batch {}, the loss is {:.2f}'.f
```

```
In [14]: # Inform that training is finished
```

```
train_finish_callback = tf.keras.callbacks.LambdaCallback(  
    on_train_end=lambda logs: print('Training finished!'))
```

```
In [15]: # Train the model with the lambda callbacks
```

```
history = model.fit(train_data, train_targets, epochs=5, batch_size=100,  
                    callbacks=[epoch_callback, batch_loss_callback,train_finish_callback])
```

Starting Epoch 1!

After batch 0, the loss is 5051.56.

After batch 1, the loss is 6418.73.

After batch 2, the loss is 6527.62.

After batch 3, the loss is 5181.47.

Starting Epoch 2!

After batch 0, the loss is 5862.76.

After batch 1, the loss is 4989.46.

After batch 2, the loss is 5366.70.

After batch 3, the loss is 7047.42.

Starting Epoch 3!

After batch 0, the loss is 5375.85.

After batch 1, the loss is 4562.59.

After batch 2, the loss is 6799.52.

After batch 3, the loss is 6536.14.

Starting Epoch 4!

After batch 0, the loss is 6156.67.

After batch 1, the loss is 6451.61.

After batch 2, the loss is 5024.42.

After batch 3, the loss is 5604.71.
Starting Epoch 5!

After batch 0, the loss is 6108.82.

After batch 1, the loss is 4620.27.

After batch 2, the loss is 6353.27.

After batch 3, the loss is 6166.57.
Training finished!

1.4 Reduce learning rate on plateau

Usage `tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=10, verbose=0, mode='auto', min_delta=0.0001, cooldown=0, min_lr=0)`

The `ReduceLROnPlateau` callback allows reduction of the learning rate when a metric has stopped improving. The arguments are similar to those used in the `EarlyStopping` callback. * The argument `monitor` is used to specify which metric to base the callback on. * The `factor` is the factor by which the learning rate decreases i.e., $\text{new_lr} = \text{factor} \times \text{old_lr}$. The `patience` is the number of epochs where there is no improvement on the monitored metric before the learning rate is reduced. * The `verbose` argument will produce progress messages when set to 1. * The `mode` determines whether the learning rate will decrease when the monitored quantity stops increasing (max) or decreasing (min). The `auto` setting causes the callback to infer the mode from the monitored quantity. * The `min_delta` is the smallest change in the monitored quantity to be deemed an improvement. * The `cooldown` is the number of epochs to wait after the learning rate is changed before the callback resumes normal operation. * The `min_lr` is a lower bound on the learning rate that the callback will produce.

Let's examine a final example.

In [16]: *# Train the model with the ReduceLROnPlateau callback*

```
history = model.fit(train_data, train_targets, epochs=100, batch_size=100,
                    callbacks=[tf.keras.callbacks.ReduceLROnPlateau(
                        monitor="loss", factor=0.2, verbose=1)], verbose=False)
```

Epoch 00013: ReduceLROnPlateau reducing learning rate to 0.00416666679084301.

Epoch 00031: ReduceLROnPlateau reducing learning rate to 0.00083333333767950535.

Epoch 00050: ReduceLROnPlateau reducing learning rate to 0.00016666667070239783.

Epoch 00060: ReduceLROnPlateau reducing learning rate to 3.333333297632635e-05.

Epoch 00076: ReduceLROnPlateau reducing learning rate to 6.666666740784422e-06.

Epoch 00088: ReduceLROnPlateau reducing learning rate to 1.3333333299669903e-06.

1.4.1 Further reading and resources

- <https://keras.io/callbacks/>
- https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/LearningRateScheduler
- https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/CSVLogger
- https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/LambdaCallback

In []: