

Coding Tutorial

July 31, 2021

```
In [37]: import tensorflow as tf
         print(tf.__version__)
```

2.0.0

1 Saving and loading models

Coding tutorials ##### Section ?? ##### Section ?? ##### Section ?? ##### Section ?? ##### Section ??

Saving and loading model weights

Load and inspect CIFAR-10 dataset The CIFAR-10 dataset consists of, in total, 60000 color images, each with one of 10 labels: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck. For an introduction and a download, see [this link](#).

```
In [3]: # Import the CIFAR-10 dataset and rescale the pixel values
```

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0

# Use smaller subset -- speeds things up
x_train = x_train[:10000]
y_train = y_train[:10000]
x_test = x_test[:1000]
y_test = y_test[:1000]
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 4s 0us/step

```
In [4]: # Plot the first 10 CIFAR-10 images
```

```
import matplotlib.pyplot as plt
```

```
fig, ax = plt.subplots(1, 10, figsize=(10, 1))
for i in range(10):
    ax[i].set_axis_off()
    ax[i].imshow(x_train[i])
```



Introduce two useful functions

In [4]: *# Introduce function to test model accuracy*

```
def get_test_accuracy(model, x_test, y_test):
    test_loss, test_acc = model.evaluate(x=x_test, y=y_test, verbose=0)
    print('accuracy: {acc:0.3f}'.format(acc=test_acc))
```

In [5]: *# Introduce function that creates a new instance of a simple CNN*

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D

def get_new_model():
    model = Sequential([
        Conv2D(filters=16, input_shape=(32, 32, 3), kernel_size=(3, 3),
            activation='relu', name='conv_1'),
        Conv2D(filters=8, kernel_size=(3, 3), activation='relu', name='conv_2'),
        MaxPooling2D(pool_size=(4, 4), name='pool_1'),
        Flatten(name='flatten'),
        Dense(units=32, activation='relu', name='dense_1'),
        Dense(units=10, activation='softmax', name='dense_2')
    ])
    model.compile(optimizer='adam',
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])
    return model
```

Create simple convolutional neural network classifier

In [9]: *# Create an instance of the model and show model summary*

```
model = get_new_model()
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 30, 30, 16)	448
conv_2 (Conv2D)	(None, 28, 28, 8)	1160
pool_1 (MaxPooling2D)	(None, 7, 7, 8)	0
flatten (Flatten)	(None, 392)	0
dense_1 (Dense)	(None, 32)	12576
dense_2 (Dense)	(None, 10)	330
Total params: 14,514		
Trainable params: 14,514		
Non-trainable params: 0		

```
In [10]: # Test accuracy of the untrained model, around 10% (random)
         get_test_accuracy(model, x_test, y_test)

accuracy: 0.100
```

Train model with checkpoints

```
In [12]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [35]: # Create Tensorflow checkpoint object
```

```
checkpoint_path = 'model_checkpoints/checkpoint'
checkpoint = ModelCheckpoint(filepath=checkpoint_path, frequency='epoch', save_weights_only=True)
```

```
In [36]: # Fit model, with simple checkpoint which saves (and overwrites) model weights every epoch
```

```
model.fit(x=x_train,
          y=y_train,
          epochs=3,
          callbacks=[checkpoint])
```

Train on 10000 samples

Epoch 1/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.6219 - accuracy: 0.4257

Epoch 00001: saving model to model_checkpoints/checkpoint

10000/10000 [=====] - 46s 5ms/sample - loss: 1.6216 - accuracy: 0.4257

```
Epoch 2/3
9984/10000 [=====>.] - ETA: 0s - loss: 1.4801 - accuracy: 0.4712
Epoch 00002: saving model to model_checkpoints/checkpoint
10000/10000 [=====] - 46s 5ms/sample - loss: 1.4798 - accuracy: 0.4712
Epoch 3/3
9984/10000 [=====>.] - ETA: 0s - loss: 1.4169 - accuracy: 0.4973 ETA: 0s
Epoch 00003: saving model to model_checkpoints/checkpoint
10000/10000 [=====] - 45s 5ms/sample - loss: 1.4168 - accuracy: 0.4973
```

```
Out[36]: <tensorflow.python.keras.callbacks.History at 0x7f330d698b00>
```

```
In [37]: # Have a look at what the checkpoint creates
```

```
! ls -lh model_checkpoints/

total 184K
-rw-r--r-- 1 jovyan users 77 Jul 30 19:01 checkpoint
-rw-r--r-- 1 jovyan users 174K Jul 30 19:01 checkpoint.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:01 checkpoint.index
```

```
In [38]: # Evaluate the performance of the trained model
```

```
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.496
```

Create new model, load weights

```
In [39]: # Create a new instance of the (initialised) model, accuracy around 10% again
```

```
model = get_new_model()
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.102
```

```
In [40]: # Load weights -- accuracy is the same as the trained model
```

```
model.load_weights(checkpoint_path)#'model_checkpoints/checkpoint/variables/variables'
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.496
```

Clear directory

```
In [33]: ! rm -r model_checkpoints
```

```
rm: cannot remove 'model_checkpoints': No such file or directory
```

Model saving criteria

Create more customised checkpoint

```
In [42]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [43]: # Create Tensorflow checkpoint object with epoch and batch details
checkpoint_5000_patch = \
    'model_checkpoints_5000/checkpoint_{epoch:02d}_{batch:04d}'

checkpoint_5000 = ModelCheckpoint(filepath=checkpoint_5000_patch,
                                  save_weights_only=True,
                                  save_freq=5000,
                                  verbose=1
                                  )
```

```
In [46]: # Create and fit model with checkpoint
```

```
model = get_new_model()
model.fit(x=x_train,
         y=y_train,
         epochs=3,
         validation_data=(x_test, y_test),
         batch_size=10,
         callbacks=[checkpoint_5000])
```

Train on 10000 samples, validate on 1000 samples

Epoch 1/3

4990/10000 [=====>...] - ETA: 30s - loss: 2.0489 - accuracy: 0.2313- ETA: 42s - 1 - E

Epoch 00001: saving model to model_checkpoints_5000/checkpoint_01_0499

9980/10000 [=====>.] - ETA: 0s - loss: 1.8774 - accuracy: 0.3037

Epoch 00001: saving model to model_checkpoints_5000/checkpoint_01_0999

10000/10000 [=====] - 62s 6ms/sample - loss: 1.8767 - accuracy: 0.3037

Epoch 2/3

4990/10000 [=====>...] - ETA: 29s - loss: 1.5617 - accuracy: 0.4339

Epoch 00002: saving model to model_checkpoints_5000/checkpoint_02_0499

9980/10000 [=====>.] - ETA: 0s - loss: 1.5402 - accuracy: 0.4440

Epoch 00002: saving model to model_checkpoints_5000/checkpoint_02_0999

10000/10000 [=====] - 61s 6ms/sample - loss: 1.5403 - accuracy: 0.4437

Epoch 3/3

4990/10000 [=====>...] - ETA: 28s - loss: 1.4145 - accuracy: 0.4810

Epoch 00003: saving model to model_checkpoints_5000/checkpoint_03_0499

9990/10000 [=====>.] - ETA: 0s - loss: 1.4229 - accuracy: 0.4839

Epoch 00003: saving model to model_checkpoints_5000/checkpoint_03_0999

10000/10000 [=====] - 59s 6ms/sample - loss: 1.4230 - accuracy: 0.4837

```
Out[46]: <tensorflow.python.keras.callbacks.History at 0x7f330c81d978>
```

```
In [47]: # Have a look at what the checkpoint creates
```

```
! ls -lh model_checkpoints_5000
```

```
total 1.1M
-rw-r--r-- 1 jovyan users 93 Jul 30 19:42 checkpoint
-rw-r--r-- 1 jovyan users 174K Jul 30 19:40 checkpoint_01_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:40 checkpoint_01_0499.index
-rw-r--r-- 1 jovyan users 174K Jul 30 19:40 checkpoint_01_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:40 checkpoint_01_0999.index
-rw-r--r-- 1 jovyan users 174K Jul 30 19:41 checkpoint_02_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:41 checkpoint_02_0499.index
-rw-r--r-- 1 jovyan users 174K Jul 30 19:41 checkpoint_02_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:41 checkpoint_02_0999.index
-rw-r--r-- 1 jovyan users 174K Jul 30 19:42 checkpoint_03_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:42 checkpoint_03_0499.index
-rw-r--r-- 1 jovyan users 174K Jul 30 19:42 checkpoint_03_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:42 checkpoint_03_0999.index
```

Work with model saving criteria

In [48]: *# Use tiny training and test set -- will overfit!*

```
x_train = x_train[:100]
y_train = y_train[:100]
x_test = x_test[:100]
y_test = y_test[:100]
```

In [57]: *# Create a new instance of untrained model*

```
model = get_new_model()
```

In [58]: *# Create Tensorflow checkpoint object which monitors the validation accuracy*

```
checkpoint_best_path = 'model_checkpoint_best/checkpoint'
checkpoint_best = ModelCheckpoint(filepath=checkpoint_best_path,
                                  save_weights_only=True,
                                  save_freq='epoch',
                                  monitor='val_accuracy',
                                  save_best_only=True,
                                  verbose=1)
```

In [59]: *# Fit the model and save only the weights with the highest validation accuracy*

```
history = model.fit(x=x_train,
                    y=y_train,
                    epochs=50,
                    validation_data=(x_test, y_test),
                    batch_size=10,
                    callbacks=[checkpoint_best],
                    verbose=0)
```

WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.iter

WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_1

WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_2
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.decay
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.learning_rate
WARNING:tensorflow:A checkpoint was restored (e.g. tf.train.Checkpoint.restore or tf.keras.Model.load_weights) but it may contain unresolved symbols.

Epoch 00001: val_accuracy improved from -inf to 0.07000, saving model to model_checkpoint_best

Epoch 00002: val_accuracy did not improve from 0.07000

Epoch 00003: val_accuracy did not improve from 0.07000

Epoch 00004: val_accuracy did not improve from 0.07000

Epoch 00005: val_accuracy improved from 0.07000 to 0.08000, saving model to model_checkpoint_best

Epoch 00006: val_accuracy did not improve from 0.08000

Epoch 00007: val_accuracy improved from 0.08000 to 0.12000, saving model to model_checkpoint_best

Epoch 00008: val_accuracy did not improve from 0.12000

Epoch 00009: val_accuracy did not improve from 0.12000

Epoch 00010: val_accuracy improved from 0.12000 to 0.13000, saving model to model_checkpoint_best

Epoch 00011: val_accuracy improved from 0.13000 to 0.14000, saving model to model_checkpoint_best

Epoch 00012: val_accuracy improved from 0.14000 to 0.16000, saving model to model_checkpoint_best

Epoch 00013: val_accuracy did not improve from 0.16000

Epoch 00014: val_accuracy did not improve from 0.16000

Epoch 00015: val_accuracy improved from 0.16000 to 0.19000, saving model to model_checkpoint_best

Epoch 00016: val_accuracy did not improve from 0.19000

Epoch 00017: val_accuracy did not improve from 0.19000

Epoch 00018: val_accuracy did not improve from 0.19000

Epoch 00019: val_accuracy did not improve from 0.19000

Epoch 00020: val_accuracy did not improve from 0.19000

Epoch 00021: val_accuracy did not improve from 0.19000

Epoch 00022: val_accuracy did not improve from 0.19000

Epoch 00023: val_accuracy did not improve from 0.19000

Epoch 00024: val_accuracy improved from 0.19000 to 0.20000, saving model to model_checkpoint_b

Epoch 00025: val_accuracy did not improve from 0.20000

Epoch 00026: val_accuracy improved from 0.20000 to 0.21000, saving model to model_checkpoint_b

Epoch 00027: val_accuracy did not improve from 0.21000

Epoch 00028: val_accuracy did not improve from 0.21000

Epoch 00029: val_accuracy did not improve from 0.21000

Epoch 00030: val_accuracy improved from 0.21000 to 0.22000, saving model to model_checkpoint_b

Epoch 00031: val_accuracy did not improve from 0.22000

Epoch 00032: val_accuracy did not improve from 0.22000

Epoch 00033: val_accuracy improved from 0.22000 to 0.23000, saving model to model_checkpoint_b

Epoch 00034: val_accuracy did not improve from 0.23000

Epoch 00035: val_accuracy did not improve from 0.23000

Epoch 00036: val_accuracy did not improve from 0.23000

Epoch 00037: val_accuracy did not improve from 0.23000

Epoch 00038: val_accuracy did not improve from 0.23000

Epoch 00039: val_accuracy did not improve from 0.23000

Epoch 00040: val_accuracy did not improve from 0.23000

Epoch 00041: val_accuracy did not improve from 0.23000

Epoch 00042: val_accuracy did not improve from 0.23000

Epoch 00043: val_accuracy did not improve from 0.23000

Epoch 00044: val_accuracy did not improve from 0.23000

Epoch 00045: val_accuracy did not improve from 0.23000

Epoch 00046: val_accuracy did not improve from 0.23000

Epoch 00047: val_accuracy did not improve from 0.23000

Epoch 00048: val_accuracy did not improve from 0.23000

Epoch 00049: val_accuracy did not improve from 0.23000


Epoch 00050: val_accuracy did not improve from 0.23000

In [60]: *# Plot training and testing curves*

```
import pandas as pd
```

```
df = pd.DataFrame(history.history)
df.plot(y=['accuracy', 'val_accuracy'])
```

Out[60]: <matplotlib.axes._subplots.AxesSubplot at 0x7f330d225400>



In [61]: *# Inspect the checkpoint directory*

```
! ls -lh model_checkpoint_best/
```

total 184K

-rw-r--r-- 1 jovyan users 77 Jul 30 19:50 checkpoint

-rw-r--r-- 1 jovyan users 174K Jul 30 19:50 checkpoint.data-00000-of-00001

```
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:50 checkpoint.index
```

```
In [65]: # Create a new model with the saved weights
new_model = get_new_model()
new_model.load_weights(checkpoint_best_path)
get_test_accuracy(new_model, x_test, y_test)
```

accuracy: 0.230

Clear directory

```
In [67]: ! rm -r model_checkpoints_5000 model_checkpoint_best
```

rm: cannot remove 'model_checkpoints_5000': No such file or directory

Saving the entire model

Create checkpoint that saves whole model, not just weights

```
In [1]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [6]: # Create Tensorflow checkpoint object
```

```
checkpoint_path = 'model_checkpoints'
checkpoints = ModelCheckpoint(filepath=checkpoint_path,
                              save_weights_only=False,
                              frequency='epoch',
                              verbose=1)
```

```
In [9]: # Create and fit model with checkpoint
```

```
model = get_new_model()
model.fit(x=x_train,
        y=y_train,
        epochs=3,
        callbacks=[checkpoints])
```

Train on 10000 samples

Epoch 1/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.9769 - accuracy: 0.2631

Epoch 00001: saving model to model_checkpoints

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/reson

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

```

INFO:tensorflow:Assets written to: model_checkpoints/assets
10000/10000 [=====] - 49s 5ms/sample - loss: 1.9765 - accuracy: 0.263
Epoch 2/3
 9984/10000 [=====>.] - ETA: 0s - loss: 1.6491 - accuracy: 0.3943
Epoch 00002: saving model to model_checkpoints
INFO:tensorflow:Assets written to: model_checkpoints/assets
10000/10000 [=====] - 45s 5ms/sample - loss: 1.6497 - accuracy: 0.393
Epoch 3/3
 9984/10000 [=====>.] - ETA: 0s - loss: 1.5379 - accuracy: 0.4366
Epoch 00003: saving model to model_checkpoints
INFO:tensorflow:Assets written to: model_checkpoints/assets
10000/10000 [=====] - 45s 4ms/sample - loss: 1.5385 - accuracy: 0.436

```

Out[9]: <tensorflow.python.keras.callbacks.History at 0x7f94ac0ad940>

Inspect what the checkpoint has created

In [10]: *# Have a look at what the checkpoint creates*

```

!ls -lh model_checkpoints/

total 312K
drwxr-xr-x 2 jovyan users 6.0K Jul 30 22:37 assets
-rw-r--r-- 1 jovyan users  77 Jul 30 19:01 checkpoint
-rw-r--r-- 1 jovyan users 174K Jul 30 19:01 checkpoint.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Jul 30 19:01 checkpoint.index
-rw-r--r-- 1 jovyan users 119K Jul 30 22:38 saved_model.pb
drwxr-xr-x 2 jovyan users 6.0K Jul 30 22:38 variables

```

In [11]: *# Enter variables directory*

```

!ls -lh model_checkpoints/variables/

total 184K
-rw-r--r-- 1 jovyan users 177K Jul 30 22:38 variables.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.1K Jul 30 22:38 variables.index

```

In [12]: *# Get the model's test accuracy*

```

get_test_accuracy(model, x_test, y_test)

accuracy: 0.464

```

Create new model from scratch

```
In [14]: # Delete model
```

```
model
```

```
-----  
NameError
```

```
Traceback (most recent call last)
```

```
<ipython-input-14-4ad52f4df476> in <module>
```

```
1 # Delete model
```

```
2
```

```
----> 3 model
```

```
NameError: name 'model' is not defined
```

```
In [15]: from tensorflow.keras.models import load_model
```

```
In [18]: # Reload model from scratch
```

```
model = load_model(checkpoint_path)
```

```
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.464
```

Use the .h5 format to save model

```
In [19]: # Save the model in .h5 format
```

```
model.save('my_model.h5')
```

```
In [20]: # Inspect .h5 file
```

```
!ls -lh my_model.h5
```

```
-rw-r--r-- 1 jovyan users 77K Jul 30 22:46 my_model.h5
```

```
In [22]: # Delete model
```

```
model
```

NameError

Traceback (most recent call last)

```
<ipython-input-22-4ad52f4df476> in <module>
    1 # Delete model
    2
----> 3 model
```

NameError: name 'model' is not defined

```
In [24]: # Reload model from scratch
```

```
model = load_model('my_model.h5')
get_test_accuracy(model, x_test, y_test)
```

accuracy: 0.464

Clear directory

```
In [25]: ! rm -r model_checkpoints
        ! rm my_model.h5
```

Loading pre-trained Keras models

Import and build Keras ResNet50 model Today we'll be using the ResNet50 model designed by a team at Microsoft Research, available through Keras applications. Please see the description on the [Keras applications page](#) for details. If you continue using it, please cite it properly! The paper it comes from is:

Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun. "Deep Residual Learning for Image Recognition", 2015.

This model takes a long time to download on the Coursera platform, so it is pre-downloaded in your workspace and saved in Keras HDF5 format. If you want to import it on your personal machine, use the following code:

```
from tensorflow.keras.applications import ResNet50
model = ResNet50(weights='imagenet')
```

In this coding tutorial, you will instead load the model directly from disk.

```
In [27]: from tensorflow.keras.models import load_model
```

```
In [28]: # Build Keras ResNet50 model
```

```
model = load_model('models/Keras_ResNet50.h5')
model.summary()
```

WARNING:tensorflow:No training configuration found in save file: the model was *not* compiled.
Model: "resnet50"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	input_1[0][0]
conv1_conv (Conv2D)	(None, 112, 112, 64)	9472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 112, 112, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 112, 112, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv2_block1_1_conv[0][0]
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_1_bn[0][0]
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36928	conv2_block1_1_relu[0][0]
conv2_block1_2_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv2_block1_2_conv[0][0]
conv2_block1_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_2_bn[0][0]
conv2_block1_0_conv (Conv2D)	(None, 56, 56, 256)	16640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block1_2_relu[0][0]
conv2_block1_0_bn (BatchNormalization)	(None, 56, 56, 256)	1024	conv2_block1_0_conv[0][0]
conv2_block1_3_bn (BatchNormalization)	(None, 56, 56, 256)	1024	conv2_block1_3_conv[0][0]
conv2_block1_add (Add)	(None, 56, 56, 256)	0	conv2_block1_0_bn[0][0] conv2_block1_3_bn[0][0]
conv2_block1_out (Activation)	(None, 56, 56, 256)	0	conv2_block1_add[0][0]
conv2_block2_1_conv (Conv2D)	(None, 56, 56, 64)	16448	conv2_block1_out[0][0]
conv2_block2_1_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv2_block2_1_conv[0][0]

conv2_block2_1_relu	(Activation	(None, 56, 56, 64)	0	conv2_block2_1_bn[0][0]
conv2_block2_2_conv	(Conv2D)	(None, 56, 56, 64)	36928	conv2_block2_1_relu[0][0]
conv2_block2_2_bn	(BatchNormali	(None, 56, 56, 64)	256	conv2_block2_2_conv[0][0]
conv2_block2_2_relu	(Activation	(None, 56, 56, 64)	0	conv2_block2_2_bn[0][0]
conv2_block2_3_conv	(Conv2D)	(None, 56, 56, 256)	16640	conv2_block2_2_relu[0][0]
conv2_block2_3_bn	(BatchNormali	(None, 56, 56, 256)	1024	conv2_block2_3_conv[0][0]
conv2_block2_add	(Add)	(None, 56, 56, 256)	0	conv2_block1_out[0][0] conv2_block2_3_bn[0][0]
conv2_block2_out	(Activation)	(None, 56, 56, 256)	0	conv2_block2_add[0][0]
conv2_block3_1_conv	(Conv2D)	(None, 56, 56, 64)	16448	conv2_block2_out[0][0]
conv2_block3_1_bn	(BatchNormali	(None, 56, 56, 64)	256	conv2_block3_1_conv[0][0]
conv2_block3_1_relu	(Activation	(None, 56, 56, 64)	0	conv2_block3_1_bn[0][0]
conv2_block3_2_conv	(Conv2D)	(None, 56, 56, 64)	36928	conv2_block3_1_relu[0][0]
conv2_block3_2_bn	(BatchNormali	(None, 56, 56, 64)	256	conv2_block3_2_conv[0][0]
conv2_block3_2_relu	(Activation	(None, 56, 56, 64)	0	conv2_block3_2_bn[0][0]
conv2_block3_3_conv	(Conv2D)	(None, 56, 56, 256)	16640	conv2_block3_2_relu[0][0]
conv2_block3_3_bn	(BatchNormali	(None, 56, 56, 256)	1024	conv2_block3_3_conv[0][0]
conv2_block3_add	(Add)	(None, 56, 56, 256)	0	conv2_block2_out[0][0] conv2_block3_3_bn[0][0]
conv2_block3_out	(Activation)	(None, 56, 56, 256)	0	conv2_block3_add[0][0]
conv3_block1_1_conv	(Conv2D)	(None, 28, 28, 128)	32896	conv2_block3_out[0][0]
conv3_block1_1_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block1_1_conv[0][0]
conv3_block1_1_relu	(Activation	(None, 28, 28, 128)	0	conv3_block1_1_bn[0][0]
conv3_block1_2_conv	(Conv2D)	(None, 28, 28, 128)	147584	conv3_block1_1_relu[0][0]
conv3_block1_2_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block1_2_conv[0][0]

conv3_block1_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_2_bn[0][0]
conv3_block1_0_conv (Conv2D)	(None, 28, 28, 512)	131584	conv2_block3_out[0][0]
conv3_block1_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block1_2_relu[0][0]
conv3_block1_0_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block1_0_conv[0][0]
conv3_block1_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block1_3_conv[0][0]
conv3_block1_add (Add)	(None, 28, 28, 512)	0	conv3_block1_0_bn[0][0] conv3_block1_3_bn[0][0]
conv3_block1_out (Activation)	(None, 28, 28, 512)	0	conv3_block1_add[0][0]
conv3_block2_1_conv (Conv2D)	(None, 28, 28, 128)	65664	conv3_block1_out[0][0]
conv3_block2_1_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block2_1_conv[0][0]
conv3_block2_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_1_bn[0][0]
conv3_block2_2_conv (Conv2D)	(None, 28, 28, 128)	147584	conv3_block2_1_relu[0][0]
conv3_block2_2_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block2_2_conv[0][0]
conv3_block2_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_2_bn[0][0]
conv3_block2_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block2_2_relu[0][0]
conv3_block2_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block2_3_conv[0][0]
conv3_block2_add (Add)	(None, 28, 28, 512)	0	conv3_block1_out[0][0] conv3_block2_3_bn[0][0]
conv3_block2_out (Activation)	(None, 28, 28, 512)	0	conv3_block2_add[0][0]
conv3_block3_1_conv (Conv2D)	(None, 28, 28, 128)	65664	conv3_block2_out[0][0]
conv3_block3_1_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block3_1_conv[0][0]
conv3_block3_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_1_bn[0][0]
conv3_block3_2_conv (Conv2D)	(None, 28, 28, 128)	147584	conv3_block3_1_relu[0][0]
conv3_block3_2_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block3_2_conv[0][0]
conv3_block3_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_2_bn[0][0]

conv3_block3_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block3_2_relu[0][0]
conv3_block3_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block3_3_conv[0][0]
conv3_block3_add (Add)	(None, 28, 28, 512)	0	conv3_block2_out[0][0] conv3_block3_3_bn[0][0]
conv3_block3_out (Activation)	(None, 28, 28, 512)	0	conv3_block3_add[0][0]
conv3_block4_1_conv (Conv2D)	(None, 28, 28, 128)	65664	conv3_block3_out[0][0]
conv3_block4_1_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block4_1_conv[0][0]
conv3_block4_1_relu (Activation	(None, 28, 28, 128)	0	conv3_block4_1_bn[0][0]
conv3_block4_2_conv (Conv2D)	(None, 28, 28, 128)	147584	conv3_block4_1_relu[0][0]
conv3_block4_2_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block4_2_conv[0][0]
conv3_block4_2_relu (Activation	(None, 28, 28, 128)	0	conv3_block4_2_bn[0][0]
conv3_block4_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block4_2_relu[0][0]
conv3_block4_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block4_3_conv[0][0]
conv3_block4_add (Add)	(None, 28, 28, 512)	0	conv3_block3_out[0][0] conv3_block4_3_bn[0][0]
conv3_block4_out (Activation)	(None, 28, 28, 512)	0	conv3_block4_add[0][0]
conv4_block1_1_conv (Conv2D)	(None, 14, 14, 256)	131328	conv3_block4_out[0][0]
conv4_block1_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block1_1_conv[0][0]
conv4_block1_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block1_1_bn[0][0]
conv4_block1_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block1_1_relu[0][0]
conv4_block1_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block1_2_conv[0][0]
conv4_block1_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block1_2_bn[0][0]
conv4_block1_0_conv (Conv2D)	(None, 14, 14, 1024)	525312	conv3_block4_out[0][0]
conv4_block1_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block1_2_relu[0][0]
conv4_block1_0_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block1_0_conv[0][0]

conv4_block1_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block1_3_conv[0][0]
conv4_block1_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_0_bn[0][0] conv4_block1_3_bn[0][0]
conv4_block1_out (Activation)	(None, 14, 14, 1024)	0	conv4_block1_add[0][0]
conv4_block2_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block1_out[0][0]
conv4_block2_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block2_1_conv[0][0]
conv4_block2_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block2_1_bn[0][0]
conv4_block2_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block2_1_relu[0][0]
conv4_block2_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block2_2_conv[0][0]
conv4_block2_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block2_2_bn[0][0]
conv4_block2_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block2_2_relu[0][0]
conv4_block2_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block2_3_conv[0][0]
conv4_block2_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_out[0][0] conv4_block2_3_bn[0][0]
conv4_block2_out (Activation)	(None, 14, 14, 1024)	0	conv4_block2_add[0][0]
conv4_block3_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block2_out[0][0]
conv4_block3_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block3_1_conv[0][0]
conv4_block3_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block3_1_bn[0][0]
conv4_block3_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block3_1_relu[0][0]
conv4_block3_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block3_2_conv[0][0]
conv4_block3_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block3_2_bn[0][0]
conv4_block3_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block3_2_relu[0][0]
conv4_block3_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block3_3_conv[0][0]
conv4_block3_add (Add)	(None, 14, 14, 1024)	0	conv4_block2_out[0][0] conv4_block3_3_bn[0][0]
conv4_block3_out (Activation)	(None, 14, 14, 1024)	0	conv4_block3_add[0][0]

conv4_block4_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block3_out[0][0]
conv4_block4_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block4_1_conv[0][0]
conv4_block4_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block4_1_bn[0][0]
conv4_block4_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block4_1_relu[0][0]
conv4_block4_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block4_2_conv[0][0]
conv4_block4_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block4_2_bn[0][0]
conv4_block4_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block4_2_relu[0][0]
conv4_block4_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block4_3_conv[0][0]
conv4_block4_add (Add)	(None, 14, 14, 1024)	0	conv4_block3_out[0][0] conv4_block4_3_bn[0][0]
conv4_block4_out (Activation)	(None, 14, 14, 1024)	0	conv4_block4_add[0][0]
conv4_block5_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block4_out[0][0]
conv4_block5_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block5_1_conv[0][0]
conv4_block5_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block5_1_bn[0][0]
conv4_block5_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block5_1_relu[0][0]
conv4_block5_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block5_2_conv[0][0]
conv4_block5_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block5_2_bn[0][0]
conv4_block5_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block5_2_relu[0][0]
conv4_block5_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block5_3_conv[0][0]
conv4_block5_add (Add)	(None, 14, 14, 1024)	0	conv4_block4_out[0][0] conv4_block5_3_bn[0][0]
conv4_block5_out (Activation)	(None, 14, 14, 1024)	0	conv4_block5_add[0][0]
conv4_block6_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block5_out[0][0]
conv4_block6_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block6_1_conv[0][0]
conv4_block6_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block6_1_bn[0][0]

conv4_block6_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block6_1_relu[0][0]
conv4_block6_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block6_2_conv[0][0]
conv4_block6_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block6_2_bn[0][0]
conv4_block6_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block6_2_relu[0][0]
conv4_block6_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block6_3_conv[0][0]
conv4_block6_add (Add)	(None, 14, 14, 1024)	0	conv4_block5_out[0][0] conv4_block6_3_bn[0][0]
conv4_block6_out (Activation)	(None, 14, 14, 1024)	0	conv4_block6_add[0][0]
conv5_block1_1_conv (Conv2D)	(None, 7, 7, 512)	524800	conv4_block6_out[0][0]
conv5_block1_1_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block1_1_conv[0][0]
conv5_block1_1_relu (Activation	(None, 7, 7, 512)	0	conv5_block1_1_bn[0][0]
conv5_block1_2_conv (Conv2D)	(None, 7, 7, 512)	2359808	conv5_block1_1_relu[0][0]
conv5_block1_2_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block1_2_conv[0][0]
conv5_block1_2_relu (Activation	(None, 7, 7, 512)	0	conv5_block1_2_bn[0][0]
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2099200	conv4_block6_out[0][0]
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block1_2_relu[0][0]
conv5_block1_0_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_0_conv[0][0]
conv5_block1_3_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_3_conv[0][0]
conv5_block1_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_0_bn[0][0] conv5_block1_3_bn[0][0]
conv5_block1_out (Activation)	(None, 7, 7, 2048)	0	conv5_block1_add[0][0]
conv5_block2_1_conv (Conv2D)	(None, 7, 7, 512)	1049088	conv5_block1_out[0][0]
conv5_block2_1_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_1_conv[0][0]
conv5_block2_1_relu (Activation	(None, 7, 7, 512)	0	conv5_block2_1_bn[0][0]
conv5_block2_2_conv (Conv2D)	(None, 7, 7, 512)	2359808	conv5_block2_1_relu[0][0]

conv5_block2_2_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_2_conv[0][0]
conv5_block2_2_relu	(Activation	(None, 7, 7, 512)	0	conv5_block2_2_bn[0][0]
conv5_block2_3_conv	(Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block2_2_relu[0][0]
conv5_block2_3_bn	(BatchNormali	(None, 7, 7, 2048)	8192	conv5_block2_3_conv[0][0]
conv5_block2_add	(Add)	(None, 7, 7, 2048)	0	conv5_block1_out[0][0] conv5_block2_3_bn[0][0]
conv5_block2_out	(Activation)	(None, 7, 7, 2048)	0	conv5_block2_add[0][0]
conv5_block3_1_conv	(Conv2D)	(None, 7, 7, 512)	1049088	conv5_block2_out[0][0]
conv5_block3_1_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_1_conv[0][0]
conv5_block3_1_relu	(Activation	(None, 7, 7, 512)	0	conv5_block3_1_bn[0][0]
conv5_block3_2_conv	(Conv2D)	(None, 7, 7, 512)	2359808	conv5_block3_1_relu[0][0]
conv5_block3_2_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_2_conv[0][0]
conv5_block3_2_relu	(Activation	(None, 7, 7, 512)	0	conv5_block3_2_bn[0][0]
conv5_block3_3_conv	(Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block3_2_relu[0][0]
conv5_block3_3_bn	(BatchNormali	(None, 7, 7, 2048)	8192	conv5_block3_3_conv[0][0]
conv5_block3_add	(Add)	(None, 7, 7, 2048)	0	conv5_block2_out[0][0] conv5_block3_3_bn[0][0]
conv5_block3_out	(Activation)	(None, 7, 7, 2048)	0	conv5_block3_add[0][0]
avg_pool	(GlobalAveragePooling2	(None, 2048)	0	conv5_block3_out[0][0]
probs	(Dense)	(None, 1000)	2049000	avg_pool[0][0]

Total params: 25,636,712

Trainable params: 25,583,592

Non-trainable params: 53,120

Import and preprocess 3 sample images

```
In [29]: # Import 3 sample ImageNet images
```

```
from tensorflow.keras.preprocessing.image import load_img

lemon_img = load_img('data/lemon.jpg', target_size=(224, 224))
viaduct_img = load_img('data/viaduct.jpg', target_size=(224, 224))
water_tower_img = load_img('data/water_tower.jpg', target_size=(224, 224))
```

Use ResNet50 model to classify images

```
In [30]: # Useful function: presents top 5 predictions and probabilities
```

```
from tensorflow.keras.preprocessing.image import img_to_array
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np
import pandas as pd

def get_top_5_predictions(img):
    x = img_to_array(img)[np.newaxis, ...]
    x = preprocess_input(x)
    preds = decode_predictions(model.predict(x), top=5)
    top_preds = pd.DataFrame(columns=['prediction', 'probability'],
                             index=np.arange(5)+1)

    for i in range(5):
        top_preds.loc[i+1, 'prediction'] = preds[0][i][1]
        top_preds.loc[i+1, 'probability'] = preds[0][i][2]
    return top_preds
```

Image 1: lemon

```
In [31]: # Display image
lemon_img
```

```
Out[31]:
```



```
In [32]: # Display top 5 predictions
         get_top_5_predictions(lemon_img)
```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_classification_2015_train.npy:
40960/35363 [=====] - 0s 0us/step

```
Out[32]: prediction probability
         1    consomme    0.227801
         2     lemon     0.221758
         3    eggnog     0.151335
         4     ladle     0.0400883
         5  spotlight    0.0291972
```

Image 2: viaduct

```
In [33]: # Display image
         viaduct_img
```

```
Out[33]:
```



```
In [34]: # Display top 5 predictions
         get_top_5_predictions(viaduct_img)
```

```
Out[34]:
```

	prediction	probability
1	vault	0.367951
2	prison	0.111742
3	viaduct	0.110738
4	window_screen	0.0882307
5	fire_screen	0.0206752

Image 3: water tower

```
In [35]: # Display image
```

```
         water_tower_img
```

```
Out[35]:
```




In [36]: # Display top 5 predictions

```
get_top_5_predictions(water_tower_img)
```

Out [36]:

	prediction	probability
1	solar_dish	0.384681
2	ladle	0.196179
3	planetarium	0.116682
4	strainer	0.04999
5	jigsaw_puzzle	0.0219473

Tensorflow Hub modules

Import and build Tensorflow Hub MobileNet v1 model Today we'll be using Google's MobileNet v1 model, available on Tensorflow Hub. Please see the description on the [Tensorflow Hub page](#) for details on it's architecture, how it's trained, and the reference. If you continue using it, please cite it properly! The paper it comes from is:

Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, Hartwig Adam: "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications", 2017.

This model takes a long time to download on the Coursera platform, so it is pre-downloaded in your workspace and saved in Tensorflow SavedModel format. If you want to import it on your personal machine, use the following code:

```
module_url = "https://tfhub.dev/google/imagenet/mobilenet_v1_050_160/classification/4"
model = Sequential([hub.KerasLayer(module_url)])
model.build(input_shape=[None, 160, 160, 3])
```

In this coding tutorial, you will instead load the model directly from disk.

```
In [40]: import tensorflow_hub as hub
         from tensorflow.keras.models import load_model
```

```
In [42]: # Build Google's Mobilenet v1 model
```

```
module = load_model('models/Tensorflow_MobileNet_v1/')
model = Sequential(hub.KerasLayer(module))
model.build(input_shape=[None, 160, 160, 3])
model.summary()
```

Model: "sequential_3"

```
-----
Layer (type)                 Output Shape              Param #
-----
keras_layer (KerasLayer)     multiple                  1343049
-----
Total params: 1,343,049
Trainable params: 0
Non-trainable params: 1,343,049
-----
```

Use MobileNet model to classify images

```
In [43]: # Import and preprocess 3 sample ImageNet images
```

```
from tensorflow.keras.preprocessing.image import load_img

lemon_img = load_img("data/lemon.jpg", target_size=(160, 160))
viaduct_img = load_img("data/viaduct.jpg", target_size=(160, 160))
water_tower_img = load_img("data/water_tower.jpg", target_size=(160, 160))
```

```
In [44]: # Read in categories text file
```

```
with open('data/imagenet_categories.txt') as txt_file:
    categories = txt_file.read().splitlines()
```

```
In [45]: # Useful function: presents top 5 predictions
```

```
import pandas as pd

def get_top_5_predictions(img):
    x = img_to_array(img)[np.newaxis, ...] / 255.0
    preds = model.predict(x)
    top_preds = pd.DataFrame(columns=['prediction'],
                              index=np.arange(5)+1)
    sorted_index = np.argsort(-preds[0])
```

```
for i in range(5):
    ith_pred = categories[sorted_index[i]]
    top_preds.loc[i+1, 'prediction'] = ith_pred

return top_preds
```

Image 1: lemon

In [46]: lemon_img

Out[46]:



In [47]: get_top_5_predictions(lemon_img)

Out[47]:

	prediction
1	shower cap
2	tray
3	candle
4	brassiere
5	African chameleon

Image 2: viaduct

In [48]: viaduct_img

Out[48]:



```
In [49]: get_top_5_predictions(viaduct_img)
```

```
Out[49]: prediction
          1      viaduct
          2      pier
          3      dam
          4      prison
          5  solar dish
```

Image 3: water tower

```
In [50]: water_tower_img
```

```
Out[50]:
```



```
In [51]: get_top_5_predictions(water_tower_img)
```

```
Out[51]:          prediction
1         solar dish
2         water tower
3  aircraft carrier
4     jigsaw puzzle
5         oxygen mask
```

```
In [ ]:
```